


# Workshop on Numerical Modeling in MHD and Plasma Physics: methods, tools, and outcomes.



## High performance collisional PIC plasma simulation with GPU-based clusters

A.V.Snytnikov,

Institute of Computational Mathematics and Mathematical Geophysics SB RAS

A.A.Romanenko

Faculty of Information Technologies, NSU

G.G.Lazareva,

Peoples' Friendship University of Russia (RUDN University), Moscow,  
Russian Federation

# Why?

Quality of Particle-In-Cell simulation crucially depends on the number of particles so it is important to be able to process as many particles as possible

The present Particle-In-Cell performance

**0.5 TFLOPS**

with GPU Tesla V100 “Volta”,  
6.4 million model particles processed in 0.006 sec.,

or

**1 Giga-particle** per second

# Main equations



$$\frac{\partial f_{i,e}}{\partial t} + \vec{v} \frac{\partial f_{i,e}}{\partial \vec{x}} + \vec{F} \frac{\partial f_{i,e}}{\partial \vec{v}} = 0$$

$$\nabla \times \vec{B} = 4\pi \vec{j} + \frac{1}{c} \frac{\partial \vec{E}}{\partial t}$$

$$\nabla \times \vec{E} = -\frac{1}{c} \frac{\partial \vec{B}}{\partial t}$$

$$\nabla \cdot \vec{E} = 4\pi \rho$$

$$\nabla \cdot \vec{B} = 0$$


$$\begin{aligned} \vec{p} &= \gamma \vec{v}, \gamma^{-1} = \sqrt{1 - v^2} \\ \vec{F} &= q_{i,e} \left( \vec{E} + \frac{1}{c} [\vec{v}, \vec{B}] \right) \\ \vec{j} &= \sum_{i,e} q_{i,e} \int f_{i,e} \vec{v} d\vec{v} \\ \rho &= \sum_{i,e} q_{i,e} \int f_{i,e} d\vec{v} \end{aligned}$$

- **Boundary conditions:** periodic

# Particle-In-Cell basics



In Particle-In-Cell method plasma is simulated by means of the model particles. Motion equations of model particles are the characteristic equations of Vlasov equation.



$$\frac{\partial p_{i,e}}{\partial t} = \kappa (E + [v, B]),$$
$$\frac{\partial r_{i,e}}{\partial t} = v_{i,e}.$$

$$p_{i,e} = \frac{v_{i,e}}{\sqrt{1 - v_{i,e}^2}}, \quad \kappa_e = -1, \quad \kappa_i = m_e/m_i.$$


- Motion equations for model particles are solved with leapfrog scheme
- Maxwell equations are solved with FDTD method
- Computational complexity: 500 floating point operations per particle
- Quality of the solution crucially depends on the number of particles



# Physical validation of the code:

## Two-stream instability simulation,

thanks to **Dr. Ekaterina Genrikh**



$$x \in [0, L], \quad y, z \in [0, n h_x]$$

$$k = 2\pi/L$$

$$W \sim e^{2\gamma t}, \quad \gamma = \frac{1}{2} \frac{\partial \ln W}{\partial t}$$

$$f(v) = \frac{1}{\Delta v \sqrt{2\pi}} \exp - \frac{(v - v_0)^2}{2\Delta v^2}$$

$n_b$  – Beam density  
 $2(\Delta v)^2$  – Beam temperature

Hydrodynamical mode

$$(k \Delta v \ll \gamma)$$

$$n_b = 2 \cdot 10^{-3}, \quad \Delta v = 0.035$$

Transition mode

$$n_b = 10^{-3}, \quad \Delta v = 0.14$$

Kynetic mode

$$(k \Delta v \gg \gamma)$$

$$n_b = 2 \cdot 10^{-4}, \quad \Delta v = 0.14$$


Computation parameters

Domain length  $L = 1.2566, 1.1424$

Spatial mesh  $100 \times 4 \times 4$

Timestep  $\tau = 0.001$


Particles per cell  $lp = 50 \dots 20000$




# Domain Decomposition




Lagrange decomposition



Эйлерова декомпозиция



Mixed Lagrange-Euler decomposition



- Particles belonging to different processors (or cores, or GPUs) are shown with different colours
- Collective MPI operations (MPI\_Allreduce are shown with large hollow arrows)
- Peer-to-peer MPI operations (send/recv) are shown with small black arrows

Y


# GPU implemenation



- The particles are stored in small array attached to cells
- Particle push is performed with shared memory
- The transit of particles from one cell to another is performed with no synchronization



# The two-level multi-flow computation




- 1) Simulation domain is decomposed ( using MPI) into subdomains, attached to different GPUs
- 2) For each GPU the particles of each cell are processed with different CUDA threads





# Weak efficiency

Lomonosov supercomputer, MSU




Mesh: 100x4x4 nodes, 6.4 million particles, MPI\_Allreduce

# Particle push



- Particle push is the most expensive part of the code (up to 90 % of the total time)
- Fortunately, particle push is accelerated the best with GPU






# Recent optimizations

- Transit of model particles from cell to cell with no synchronization using buffers
- Get rid of buffers using particle labels
- Get rid of atomic operations when computing current within a cell

# Transit of model particles from cell to cell with no synchronization using buffers



# Get rid of buffers using particle labels



- Instead of moving into buffer the particle is given a label (“which direction it goes”)
- Then neighbour cells read their particles from the particle list of the present cells with no read-and write conflict
- After that the flow-away particles are removed from the list by a separate kernel
- In such a way, send buffers (quite big) are not necessary



# Get rid of atomic operations when computing current within a cell



The particles are stored in small array attached to cells


Current is computed in each cell asynchronously by CUDA threads, one thread for one particle

Each thread **reads** the present value of the current in a mesh node, adds the current of the particle and **writes** the result into the mesh node (reading and writing to different places by different threads at once)

In this case the read-and write conflict occurs and atomic operations are necessary

If we do the opposite and attach a CUDA thread to the mesh nodes instead of particles,

Then the thread reads the current value from each particle and writes to its own memory (reading from different places but writing to one place only for each thread)



Arrows of different colours do not cross in one node.  
No read and write conflict

